



seL4 and CAmkES Real Time

Ihor Kuz, Stephen Sherratt, Hesham Almatary

August 2016

www.csiro.au



Overview




- seL4-RT: Real Time support in seL4
- CAmkES-RT: Real Time support in CAmkES
- Exercises

seL4-RT: Real Time Support in seL4

Legacy seL4 Scheduling Model

- Thread represented by TCB
- TCB has: CSpace, VSpace, IPC buffer
- Preemptive, Round Robin, 256 priority levels



Legacy seL4 Thread Creation



1. Create TCB Object

- `seL4_Untyped_Retype()`

2. Set Cspace, Vspace, IPC buffer, priority

- `seL4_TCB_Configure()`

3. Set stack and instruction pointer

- `seL4_TCB_WriteRegisters()`

4. Activate thread

- `seL4_TCB_Resume()`

New seL4 Real Time Support



Everything stays the same except...

seL4 Thread Attributes

- CSpace, VSpace
- IPC buffer
- Priority
- Timeslice

seL4-RT Thread Attributes

- CSpace, VSpace
- IPC buffer
- Priority
- Scheduling Context



Not Runnable
if NULL

Scheduling Context



- New kernel-object type
 - Bound to a TCB
 - Specifies how long a thread can be run for
- Consists of a **budget (b)** and a **period (T)** ($b \leq T$)
 - Thread consumes budget and is preempted when it exhausts the budget
 - Budget is refreshed every period
 - *kernel will not permit a thread to run for more than b out of T microseconds*
- (b, T) forms an upper bound of time usage, *not* reservation
 - ... but supports schedulability analysis that can guarantee lower bounds
 - needs to take priorities into account

Schedulability




- A thread of priority P will meet its deadline if it passes RMS schedulability test
- Utilisation:

$$u = \sum_{i=1}^n \frac{b_i}{T_i} \leq n \times \left(2^{\frac{1}{n}} - 1\right)$$

- (where priority of $T_i \leq P$)

It is the user's responsibility to make sure the system is schedulable!

Example: seL4-RT Scheduling



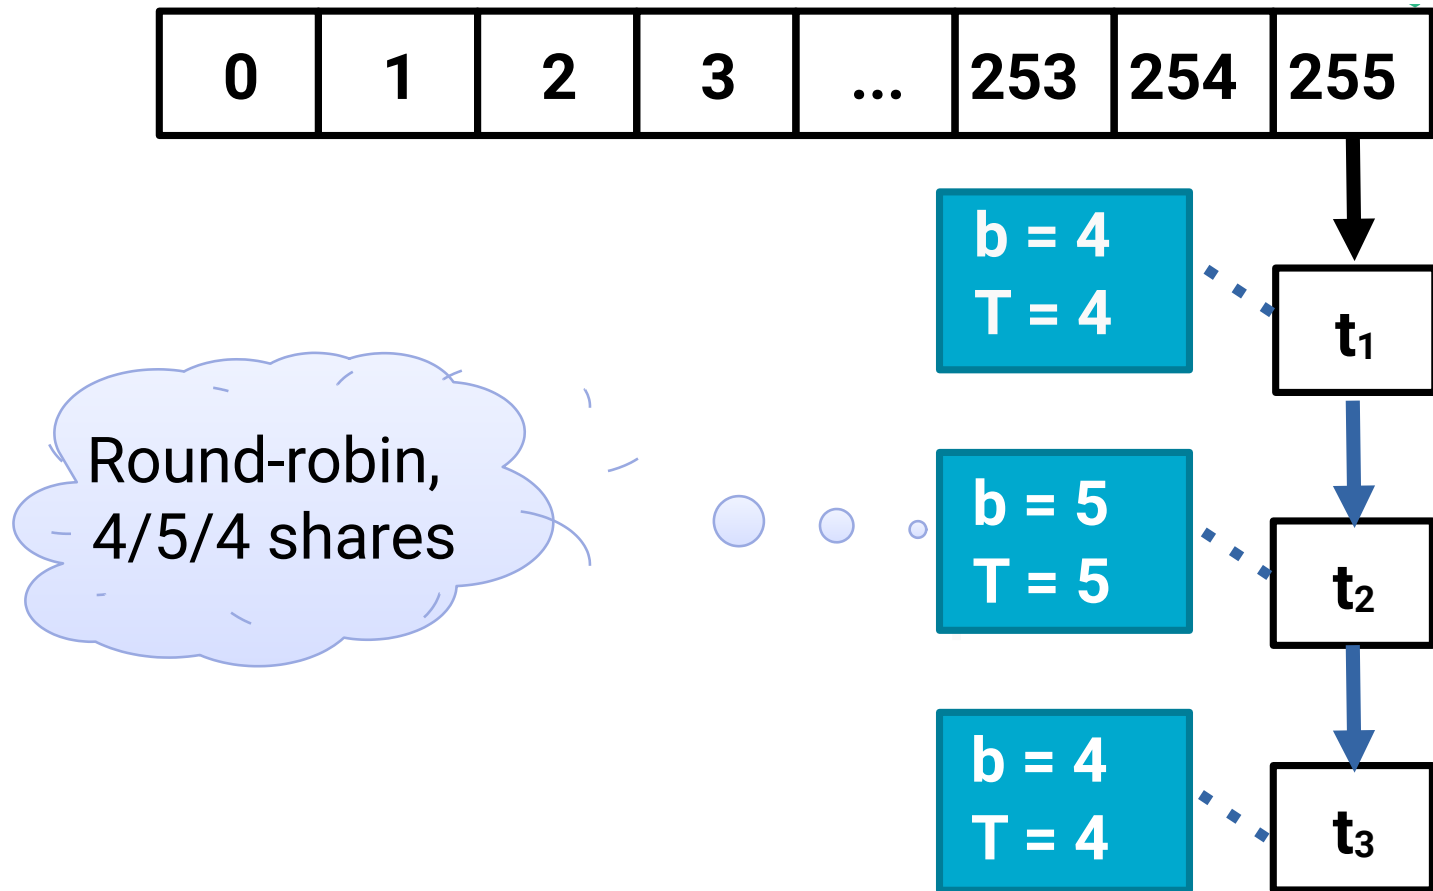
- Utilisation:

$$u = \frac{1}{8} + \frac{2}{5} + \frac{1}{5} = 0.775$$

$$n \times (2^{\frac{1}{n}} - 1) = 3 \times (2^{\frac{1}{3}} - 1) = 0.77976$$

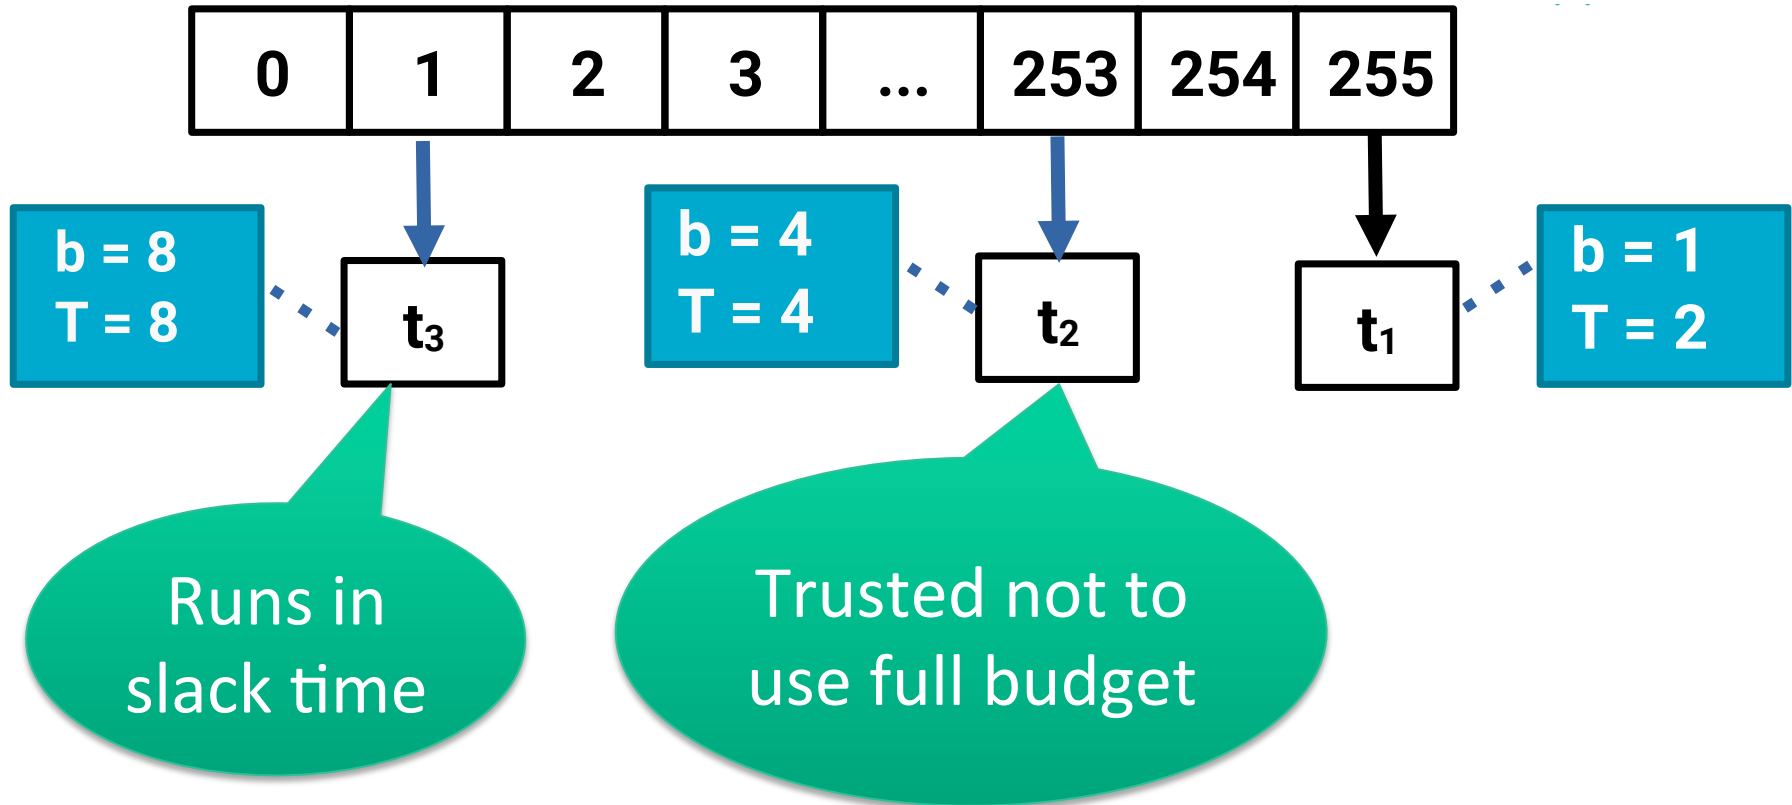
$$0.775 \leq 0.77976$$

Overcommitting: Full Budget



- utilisation: $u = 4/4 + 5/5 + 4/4 = 300\%$
- emulates legacy model where period corresponds to time slice

General Budget



seL4-RT Thread Creation



1. Create TCB Object
 - `seL4_Untyped_Retype()`
2. Set Cspace, Vspace, IPC buffer, priority
 - `seL4_TCB_Configure()`
3. Set stack and instruction pointer
 - `seL4_TCB_WriteRegisters()`
4. Create Scheduling Context Object
 - `seL4_Untyped_Retype()`
5. Assign Budget and Period
 - `seL4_SchedControl_Configure()`
6. Bind Thread to Scheduling Context
 - `seL4_SchedContext_bindTCB()`
7. Activate thread
 - `seL4_TCB_Resume()`


Passive Thread




- Thread without a scheduling context
- Use `seL4_SchedContext_UnbindTCB()` to remove scheduling context
- Different from suspending a thread
 - Suspended thread will be run again if it is the next runnable thread
- Cannot execute without the action of another thread
 - To give it a scheduling context again

A passive thread will NOT be schedulable again until it receives a scheduling context

Scheduling Context donation



Scheduling Context donation



Scheduling Context Donation



- Implicit donation via `seL4_Call()` & `seL4_NBSendRecv()`
 - ... iff the receiver does not currently have a scheduling context
 - Calling thread's scheduling context will be donated
 - Calling thread will no longer be schedulable
- Scheduling context returned when `seL4_Reply()` or `seL4_ReplyRecv()` is invoked
 - Scheduling context might *NOT* return if passive thread doesn't reply
- Use `seL4_SchedContext_BindNotification()` to donate via notification object

New APIs



- `seL4_SchedControl_Configure()`
 - Invoke SchedControl object to set budget and period of a SchedContext
- `seL4_SchedContext_Yield()`
 - Give up the rest of the budget until the next budget refresh
- `seL4_SchedContext_BindTCB()`
 - Give a TCB a SchedContext
- `seL4_SchedContext_UnbindTCB()`
 - Take a SchedContext away from a TCB
- `seL4_SchedContext_BindNotification()`
 - Allow a SchedContext to be donated via a Notification
- `seL4_SchedContext_UnbindNotification()`
 - Stop donation via a Notification

Warning: Priority Type Change



seL4 priority

- `uint8_t` priority
- Thread runs at a fixed priority

seL4-RT priority

- `seL4_Prio_t` priority
 - `uint8_t` priority
 - `uint8_t` `max_priority`
- Thread can change (own or other's) priority
 - ceiling defined by `max_priority` value

seL4-RT Exercises



Real Time Support in CAmkES

Threads in CAmkES



- Control thread (`run()` function)
- All interfaces run in separate threads

How CAmkES Supports seL4-RT



New pre-defined attributes

- Scheduling context (sc) attribute
 - By default each thread has its own scheduling context
 - Passive thread: explicitly sets scheduling context to “none”
- Budget attribute
 - Set a thread’s default scheduling context’s budget
- Period attribute
 - Set a thread’s default scheduling context’s period
- Defaults
 - Period = budget = 10,000 (microseconds)

Using Real Time Attributes



Configure scheduling context for:

- All threads in one component
 - `$(instance).budget`
 - `$(instance).period`
- Control thread
 - `$(instance)._budget`
 - `$(instance)._period`
- Specific interface thread
 - `$(instance).$(interface)_budget`
 - `$(instance).$(interface)_period`
- Scheduling context attribute
 - `$(instance).$(interface)_sc`

```
Component Server {  
    provides Simple a;  
    provides Simple b;  
}  
  
Assembly {  
    composition {  
        component Server s;  
        component Client c;  
    }  
    configuration {  
        c.budget = 1000;  
        c._period = 5000;  
        s.a_period = 5000;  
        s.b_sc = "none";  
    }  
}
```

Porting Existing Apps to CAmkES-RT



- All threads in a component instance need a budget and a period
 - except passive components.
 - The default budget and period are 10000 μ s
- Passive component only make sense for the "receiving" side of RPC or Event interface
 - Don't make control thread passive!
- Always use interface-specific attribute to configure passive component
- Passive component only works with seL4RPCCall and seL4Notification connectors
 - Will do scheduling context donation

CAmkES-RT Exercises: server-1



- Get RT tutorial exercises
 - `mkdir sel4-rt-tutorials; cd sel4-rt-tutorials`
 - `repo init -u https://github.com/sel4proj/sel4-tutorials-manifest.git -m sel4-rt-tutorials.xml`
 - `repo sync`
- Server-1
 - Goal: create and initialise a passive thread, do scheduling context donation
- TODOs:
 - TODO 1: initialise the new TCB
 - TODO 2: Create a cap to store a sched context
 - TODO 3: initialise a scheduling context
 - TODO 4: Bind the scheduling context
 - TODO 5: Make this server passive
 - TODO 6: Get signal from server saying it initialised
 - TODO 7: Unbind the server's scheduling context

CAmkES-RT Exercises: server-2



- Server-2
 - Goal: experiment with budget and period settings for an active thread
- TODOs
 - TODO_4: Play around with the parameters and run the app. Try and see if you can explain what happens
- CAmkES RT Exercises
 - camkes-rt-1
 - Make a CAmkES app that does what server-1 does
 - camkes-rt-2
 - Make a CAmkES app that does what server-2 does

More info



- <http://sel4.systems/Info/Docs/seL4-manual-1.0.0-rt-dev.pdf>
 - Documentation of seL4-RT kernel
- <https://github.com/seL4/seL4/tree/rt>
 - “rt” branch of kernel
- <https://github.com/seL4/sel4test/tree/rt>
 - “rt” branch of sel4test
- <https://github.com/seL4/camkes-tool>
 - CAMkES-RT support is in camkes master branch
- <https://github.com/seL4-projects/camkes-sc-tests>
 - Examples and tests for CAMkES-RT